

Elsevier Editorial System(tm) for Energy  
Manuscript Draft

Manuscript Number:

Title: Scalable Tuning of Building Models to Hourly Data

Article Type: Full Length Article

Keywords: Autotune; EnergyPlus; calibration; optimization; evolutionary computation; automated measurement and verification

Corresponding Author: Dr. Joshua New, Ph.D.

Corresponding Author's Institution: Oak Ridge National Laboratory

First Author: Joshua New, Ph.D.

Order of Authors: Joshua New, Ph.D.; Aaron Garrett, Ph.D.

Abstract: Energy models of existing buildings are unreliable unless calibrated so they correlate well with actual energy usage. Manual tuning requires a skilled professional, is prohibitively expensive for small projects, imperfect, non-repeatable, non-transferable, and not scalable to the dozens of sensor channels that smart meters, smart appliances, and cheap/ubiquitous sensors are beginning to make available today. A scalable, automated methodology is needed to quickly and intelligently calibrate building energy models to all available data, increase the usefulness of those models, and facilitate speed-and-scale penetration of simulation-based capabilities into the marketplace for actualized energy savings. The "Autotune" project is a novel, model-agnostic methodology which leverages supercomputing, large simulation ensembles, and big data mining with multiple machine learning algorithms to allow automatic calibration of simulations that match measured experimental data in a way that is deployable on commodity hardware. This paper shares several methodologies employed to reduce the combinatorial complexity to a computationally tractable search problem for hundreds of input parameters. Accuracy metrics are provided which quantify model error to measured data for either monthly or hourly electrical usage from a highly-instrumented, emulated-occupancy research home.

**Energy Journal Cover Letter**

Joshua New  
Oak Ridge National Laboratory  
P.O. Box 2008, MS 6324  
Oak Ridge, TN 37831  
[newjr@ornl.gov](mailto:newjr@ornl.gov)

June 17, 2014

Dear Dr. H. Lund,

We wish to submit a new manuscript entitled “Scalable Tuning of Building Models to Hourly Data” for consideration by the *Energy Journal*. We confirm that this work is original and has not been published elsewhere nor is it currently under consideration for publication elsewhere.

In this paper, we report on techniques to address computational tractability for performing automated measurement and verification that calibrates a software model to measured data. We have developed a scalable method, based on evolutionary computation and multi-parameter optimization, to *automatically* calibrate dozens of input parameters for a building model such that simulation output matches hourly measured energy consumption data from a heavily instrumented, automated-occupancy research home. This is significant because many countries have legal requirements on what constitutes a “legally usable” model, such as ASHRAE Guideline 14 or the International Performance Measure & Verification Protocol, which quantitatively define the criteria for a calibrated software model so that it can be used for building codes, tax tradeoffs, financing, etc. However, the cost for creating a calibrated model currently prohibits the use of software models to the largest building projects, significantly curtailing the impact that optimized energy savings could have in light commercial or residential buildings.

This work is directly relevant to the Energy Journal as the methodology is domain-agnostic; it has been optimized for building energy consumption, but the techniques can also be used in any area where an energy model needs to be calibrated to measured data. In the U.S., buildings are the largest energy-consuming sector at 41% of primary energy production, split evenly between ~4.9 million commercial buildings and ~114 million residential homes. This work aligns with many of Energy’s research topics. Energy modelling, prediction from a calibrated model (e.g. optimal retrofits for energy efficiency), model-in-the-loop integrated energy systems for smart control, and simulation-informed energy planning/management are applicable with different scales as to the number and fidelity of calibrated software models. The paper should be of interest to readers in the areas of building science, evolutionary computation, and multi-parameter optimization.

Please address all correspondence concerning this manuscript to me at [newjr@ornl.gov](mailto:newjr@ornl.gov). Thank you for your consideration of this manuscript and I stand by to assist with any further information required.

Sincerely,

Joshua New  
R&D Staff, Oak Ridge National Laboratory  
Joint Faculty, The University of Tennessee

### **Energy Journal Highlights**

- We survey guidelines and calibration requirements for energy models of buildings.
- We apply evolutionary computation for calibrating software to measured data.
- We showcase abbreviated schedules for speeding up autonomous calibration.
- We demonstrate an island-model for parallelizing multi-parameter optimization.
- We analyze automatic calibration to hourly data with comparison to manual effort.

1  
2  
3  
4  
5  
6  
7  
8  
9

# Scalable Tuning of Building Models to Hourly Data

10  
11  
12 Aaron Garrett<sup>a</sup>, Joshua New<sup>b,\*</sup>  
13

14 <sup>a</sup>*Mathematical, Computing, and Information Sciences, Jacksonville State University,  
15 Jacksonville, AL, USA*

16 <sup>b</sup>*Oak Ridge National Laboratory, Oak Ridge, TN, USA*

17  
18  
19  
20

---

## Abstract

21  
22  
23 Energy models of existing buildings are unreliable unless calibrated so that  
24 they correlate well with actual energy usage. Manual tuning requires a skilled  
25 professional; and it is prohibitively expensive for small projects, imperfect,  
26 non-repeatable, non-transferable, and not scalable to the dozens of sensor  
27 channels that smart meters, smart appliances, and cheap, ubiquitous sen-  
28 sors are beginning to make available. A scalable, automated methodology  
29 is needed to quickly and intelligently calibrate building energy models to all  
30 available data, increase the usefulness of those models, and facilitate speed-  
31 and-scale penetration of simulation-based capabilities into the marketplace  
32 for actualized energy savings. The “Autotune” project is a novel, model-  
33 agnostic methodology that leverages supercomputing, large simulation en-  
34sembles, and big data mining with multiple machine learning algorithms to  
35 allow automatic calibration of simulations that match measured experimental  
36 data in a way that is deployable on commodity hardware. This paper shares  
37 several methodologies employed to reduce the combinatorial complexity to a  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53

---

54 \*Corresponding author at: Oak Ridge National Laboratory; Email address:  
55 newjr@ornl.gov

56  
57  
58  
59  
60  
61  
62  
63  
64  
65

1  
2  
3  
4  
5  
6  
7  
8  
9 computationally tractable search problem for hundreds of input parameters.  
10 Accuracy metrics are provided that quantify model error to measured data  
11 for either monthly or hourly electrical usage from a highly instrumented,  
12 emulated-occupancy research home.  
13

14  
15 *Keywords:* Autotune, EnergyPlus, calibration, optimization, evolutionary  
16 computation  
17  
18

---

19  
20  
21 **1. Introduction**  
22  
23

24 Sustainability, with its connections to energy and climate change, is per-  
25 haps the defining challenge of our time. With only 4.4% of the world's popu-  
26 lation, the United States (US) consumes 19% of the world's primary energy  
27 production. The largest sector of energy consumption in the US is buildings,  
28 which consumed 41% of national primary energy production in 2010; 19%  
29 was consumed by approximately 4.7 million commercial buildings and 22%  
30 by 114 million residential buildings [1]. Building energy model creation and  
31 simulation can provide many functions, but is often fiscally infeasible for all  
32 but the largest of projects because of the time required to create a model of  
33 an existing building and calibrate it to measured data. The US Department  
34 of Energy's (DOE) Building Technologies Office (BTO) is assisting devel-  
35 opment of several Emerging Technology applications to significantly reduce  
36 costs and drive simulation-informed actualized energy savings into existing  
37 light commercial and residential buildings to meet the nations overarching  
38 goal to reduce building energy use by 50% by 2030 compared with a 2010  
39 baseline.  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54

55 Many simulation-based analysis tools are available which are used to  
56  
57  
58

project how specific policies or retrofit packages would maximize return-on-investment for subsidies through federal, state, local, and utility tax incentives, rebates, and loan programs. Resolving well-known issues such as principle-agent, first cost, and cost/performance trade-offs, as well as maximizing financial metrics such as net present value, rate of return, cost of conserved energy, and simple pay-back, fall within the domain of these tools. As is the case for all software tools, their analysis suffers from “garbage in, garbage out” (GIGO). This is complicated by the fact that, unlike cars or planes built to a strict engineering specification, buildings are currently based on one-off designs and constructed in the field. These buildings can last decades or hundreds of years, and rarely do any sensor data exist beyond utility bill data. For older buildings, optimal retrofit packages and similar analyses are calculated for a fictitious building and necessarily yield suboptimal results. A central challenge in the domain of building energy efficiency is being able to realistically and cost-effectively model existing buildings. Even coarse models are useful to determine incremental energy conservation measure changes to/impact upon whole-building energy consumption. That usefulness is dramatically increased for existing buildings, for which existing data can be used to calibrate the energy model. However, differences between models and actual monthly utility bills on the order of 24–97% [2, 3] are currently common. Many measurement and verification (M&V) protocols specify a required accuracy for a model to be legally useful. Most large organizations use ASHRAE Guideline 14, 5.3.2.4.f requirements, which specify a coefficient of variance for root mean squared error (RMSE) <15% or 30% and a normalized mean bias error <5% or 10% for calibrating to monthly or

1  
2  
3  
4  
5  
6  
7  
8  
9 hourly data, respectively [4].  
10

11 Several simulation engines, and tools that leverage the simulation en-  
12 gines, are actively supported by DOE [5]. DOEs flagship simulation engine  
13 is currently EnergyPlus [6], which has been supported with over \$65 million  
14 since 1995. OpenStudio [7] now serves as the primary middleware between  
15 simulation engines and analysis tool applications. High-level graphical inter-  
16 faces and low-level, text-based files allow a user to provide information that  
17 fully describes a given building and from which EnergyPlus can calculate de-  
18 tailed heat flow and energy usage information for the building. The number  
19 and instances of these input parameters are extensive and highly variable in  
20 their combinatorial effects, and their sensitivities are not yet fully explored.  
21 This relegates simulation calibration to an “art,” and only a few hundred  
22 people have qualified for ASHRAE’s building energy modeling professional  
23 certification. It is unrealistic to expect even an advanced user to be capable  
24 of providing accurate values for each of the approximately 3,000 parameters  
25 expected by EnergyPlus for the average building. To mitigate such issues,  
26 a user often uses a reference or template building already in the preferred  
27 tool, which is similar to the user’s own building, as a default point for pa-  
28 rameter values. These values are then “corrected” to more closely match the  
29 actual building under consideration, depending on the level of information  
30 available, such as the data specified by an ASHRAE level 1, 2, or 3 audit. In  
31 addition, average material properties are typically used from the *ASHRAE*  
32 *Handbook of Fundamentals* (HoF) which is now beginning to include signif-  
33 icant variances in material properties identified from controlled laboratory  
34 tests [8].  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

1  
2  
3  
4  
5  
6  
7  
8  
9 As the variance in material properties increases; as the different systems,  
10 envelopes, equipment, components, and materials become more complicated  
11 and diverse; and as energy simulation modeling algorithms evolve to more  
12 thoroughly model existing systems and capture new equipment technologies,  
13 there will be a growing need to mitigate this increase in complexity by relying  
14 on cost-effective, intelligent algorithms to calibrate building energy models  
15 in a way that uses as much data as is available. The Autotune project [9]  
16 aims to solve this need with an automated process and has previously demon-  
17 strated calibration results for envelope parameters using monthly utility data  
18 [10]. This paper extends that work by discussing the scalable methodologies  
19 used to tune a building energy model's 100+ envelope parameters to whole-  
20 building hourly electrical usage data.  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34

## 2. Background

35 This section discusses a quick survey of Autotune project results, pre-  
36 vious surveys of simulation accuracy for multiple simulation engines, the  
37 causes for this inaccuracy, background on the process of calibrating energy  
38 models to measured data, the computational complexity involved in solving  
39 the calibration problem, and considerations in the selection of evolutionary  
40 computation methodologies employed for automated calibration.

### 41 2.1. Autotune Background

42 The Autotune project has used 269+ channels of 15-minute sensor data  
43 from a robotically emulated-occupancy ZEBRAAlliance [11, 12] 2800 ft<sup>2</sup> re-  
44 search home. Parametric ensemble models of this building were simulated  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

1  
2  
3  
4  
5  
6  
7  
8  
9 using high performance computing (HPC) resources. The Titan supercom-  
10 puter at Oak Ridge National Laboratory ORNL), currently the worlds second  
11 fastest supercomputer, allowed the use of 131,072 cores to calculate 524,288  
12 simulations and write 44 TB of data to disk in 68 minutes [13]. Some of the  
13 latest advances in web-oriented database storage were used to allow queryable  
14 simulations generated from varying 156 inputs and reporting 96 outputs at  
15 15-minute resolution (35 MB per simulation) for 8 million EnergyPlus simula-  
16 tions [14]. Measured data often are corrupt because of uncalibrated sensors or  
17 missing data, so statistical techniques have been refined for autonomous qual-  
18 ity assurance and gap-filling [15]. Extensive big data mining was conducted  
19 through creation of an HPC-enabled suite of machine learning algorithms  
20 (MLSuite [16]) to generate agent-based encapsulation of knowledge for Au-  
21 totune deployment on mobile devices. EnergyPlus was approximated with  
22 machine learning algorithms to reduce simulation runtime from 3 minutes to  
23 4 seconds with minimal trade-off in accuracy for the processed building types  
24 [16]. The Autotune project, in an effort to promote open science, is making  
25 a portion of the 267 TB (26.9 trillion data points) of EnergyPlus simulation  
26 data freely available online<sup>1</sup>.  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

## 44     *2.2. Simulation Accuracy*

45  
46  
47 Despite the proliferating use of building energy tools via their integration  
48 with research and development (R&D) initiatives, policies, tax structure,  
49 codes and standards, and product sales, there remain many concerns and  
50 shortcomings applicable to all simulation engines. The primary concern is  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

---

<sup>1</sup><http://autotune.roofcalc.com>

1  
2  
3  
4  
5  
6  
7  
8  
9 typically the accuracy of the simulation engines for realistically modeling (via  
10 inputs) a virtual building so that it matches a real-world building. This is a  
11 well-established issue that has remained unresolved in the energy modeling  
12 community. A Home Energy Rating System (HERS) study in 1999 [17]  
13 using the REM/Rate simulation engine for 2,300 homes in Wisconsin found  
14 that the median home's heating use, which accounts for 40% of the average  
15 annual Wisconsin energy bill, was overestimated by 22%, with the worst 15%  
16 median being off by 62%. A 2008 pilot study [2] found 190 Home Energy  
17 Saver, REM/Rate, and SIMPLE residential simulation models had a 25.1–  
18 96.6% error rate compared with actual monthly electrical energy usage. A  
19 recent 2012 study [3] found that 859 residential models across Home Energy  
20 Saver, REM/Rate, and SIMPLE simulation engines had a mean absolute  
21 percent difference of 24% compared with actual monthly electrical energy  
22 usage and of 24–37% compared with actual natural gas use for a sample size  
23 of 500 houses. Note that all of these studies use comparisons with monthly  
24 utility bill data; the challenge of accurately matching hourly or 15-minute  
25 data for dozens of sub-metered data channels is significantly more difficult.  
26  
27

28 The challenge for simulation accuracy can be reduced to two primary  
29 issues: (1) a gap between the as-modeled and as-built structure and (2) lim-  
30 itations of the modeling engine's capabilities. We will discuss each of these  
31 in detail.

### 32 2.3. Common Errors with Simulation Inputs

33 Gaps between as-modeled and as-built structures have many sources, with  
34 the fault being traceable to an inaccurate input file rather than the simula-  
35 tion engine itself. We have worked with building scientists and conducted a  
36

1  
2  
3  
4  
5  
6  
7  
8  
9 sensitivity analysis to identify the most important input parameters.  
10

11 Infiltration—the rate at which air and the energy in it flows through the  
12 building envelope (typically measured in cubic feet per minute per square  
13 ft)—currently cannot be cheaply tested despite being one of the most impor-  
14 tant factors for building energy efficiency. Blower-door tests can determine  
15 the infiltration rate at a given pressure (usually 50 Pascals), but these are one-  
16 time measurements that, in reality, vary significantly as a function of other  
17 variables such as temperature, wind speed, and wind direction. Therefore,  
18 infiltration is often one of the first variables most energy modeling experts  
19 use to manually align a simulation model with actual data.  
20  
21

22 A second issue is the schedule of building usage, which includes the  
23 number of occupants; times of occupancy; heating, ventilation, and air-  
24 conditioning (HVAC) set points; operation schedule; and many other fac-  
25 tors. These also constitute inputs to the simulation engine but are often  
26 specified in a separate EnergyPlus file for convenience. For many of these,  
27 cost-effective sensors simply do not exist or are not typically deployed in a  
28 building (especially sensors providing data that can be easily leveraged by  
29 energy modelers). In many cases, estimates of occupancy schedules and rel-  
30 atively static set point temperatures are estimated and then used later to  
31 “true-up” the simulation to match whole-building data without regard to  
32 the accuracy of the actual HVAC thermostat set points.  
33  
34

35 A third component that modelers often do not use, but which laboratory  
36 testing has exposed as a major contributor to model variance, is material  
37 properties used within a building model. Modelers typically use material  
38 properties from the ASHRAE HoF, which relates the average value of a  
39  
40

physical property (e.g., roughness, conductivity, density, specific heat, thermal absorptance) for a given material (e.g., 5/8 in. gypsum board, cultured stone, 2<sup>3</sup>/8 in. glass fiber insulation, 1/2 in. plywood, 2x6 in. wood studs) based on standard tests. Occasionally, materials ship from a manufacturer with labels that give more reliable estimates than the average values found in the HoF. However, laboratory-controlled hotbox testing of specific materials has shown significant variance in materials even from a single manufacturer. Although many of these values are used to update the HoF's numbers over time, data have traditionally not been reported for variance. However, this is expected to change in future versions of the HoF. In such cases, energy modelers have very few reliable data to determine the precise values necessary to create an accurate model of material properties for the construction of an individual building.

A fourth gap is that energy modelers typically use building design, building information model (BIM) files (increasingly), or similar documents for energy modeling; but craftsmanship influences how those designs are implemented (e.g., a contractor might neglect to put insulation in a corner wall). Many other gaps exist, and it is important to note that most of these gaps remain because current business models find that saving relatively cheap energy is not an adequate incentive to justify the energy modeling expense. For example, a financial analysis of 26 Federal Energy Management Program (FEMP) energy service company (ESCO) projects on commercial buildings showed that project development costs involving simulation ranged from 10–45% of the total cost for projects smaller than \$1 million. This increased risk prohibits financing for traditional, simulation-informed calibration and

1  
2  
3  
4  
5  
6  
7  
8  
9 optimization of building retrofits for all but the largest building renovation  
10 projects.  
11  
12

13  
14 *2.4. Common Errors with Simulation Algorithms*  
15

16 The limitations of engine modeling capabilities are a well-understood and  
17 active issue for funded development teams or active communities/ecosystems  
18 behind the most popular simulation engines and tools. It is common for there  
19 to be a backlog of hundreds of user requests involving the inaccuracies of  
20 specific algorithms. Unlike the one-off use of a calibrated simulation model  
21 for a single building, significant market incentives and policy impacts would  
22 result from developing a capable software engine that could be applied to the  
23 entire US building stock. Simulation engines are often used to inform policy  
24 decisions toward cost-effectively meeting national energy goals. Although  
25 several companies have developed simulation engines that are closely tied to  
26 their product lines, an increasing number are either using co-simulation or  
27 switching to the large simulation engines like EnergyPlus. However, there  
28 are a few primary factors that may be considered regarding the inaccuracies  
29 of simulation engines.  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42

43 First, most simulation engines are engineering models that attempt to  
44 model, through time and with some degree of fidelity, the underlying physics  
45 involved in energy consumption. As is necessary for such an approach,  
46 engineering algorithms are an approximation of reality (e.g., they use 1-  
47 dimensional heat transfer processes because 3-dimensional heat transfer would  
48 take too long to calculate). Some simulation engines allow users to select the  
49 needed level of precision on an algorithmic level to allow customization for a  
50 particular area of interest.  
51  
52  
53  
54  
55  
56  
57  
58

1  
2  
3  
4  
5  
6  
7  
8  
9     Statistical models show some promise, but they are more useful for normative-  
10 based models for policy decisions than for product-level or system-level mod-  
11 ifications to a building. Second, there is a lag between the development of  
12 new, innovative technologies and the capability to accurately model them  
13 within a simulation engine. Only the most active simulation engine de-  
14 velopment teams are able to keep up with, or foresee the need to model,  
15 building products, components, or systems before they achieve a significant  
16 market share. As the code base grows, the challenge of maintaining a soft-  
17 ware architecture that can accommodate new, integrated technologies (which  
18 may impact several parts of a building) grows in like manner. Third, vir-  
19 tually all known simulation engines are single-threaded; only recently have  
20 attempts been made to leverage traditional multi-core and graphical process-  
21 ing unit (GPU) computational hardware. Given the significantly different  
22 multi-threaded software development paradigms, there will be substantial  
23 challenges in scaling additional simulation capabilities without an increase  
24 in runtime for models that use those capabilities. Fourth, the computer itself  
25 is an approximation engine and is fundamentally limited with respect to the  
26 accuracy that it can provide in a unit of time for a given algorithm. For  
27 most simulation engine development, the focus has been a reactive process  
28 of building a tool that is sufficient to meet some small fraction of the long  
29 list of needs expressed by the users.

30  
31     The Autotune project [9] goal is to create an automated process for tuning  
32 simulation inputs so that simulation output matches measured data. This  
33 repeatable, transferable, scientifically rigorous process can address inaccura-  
34 cies in both input and algorithms by (1) adapting a model to more closely  
35

match real-world data from the as-built structure and (2) doing so in a way that accommodates inaccuracies in the underlying engineering model.

### 2.5. Tuning of Building Models

While inaccurate models can be useful in comparative analysis, calibrated models are more useful and often must meet specific calibration criteria to be legally permitted for use in a specific context. To satisfy these requirements, modification of simulation algorithms to address address the mismatch (often by working around limitations of a given simulation engine) remains one of the most tractable methods used by experts. As mentioned earlier, this tuning process remains an “art” that even the practitioners usually do not enjoy. Informal interviews and surveys have indicated that less than 3% of people actually enjoy the tuning process; instead, they see it as a necessary and laborious process. These practitioners often indicate that they use infiltration and schedules (of many types) as the primary “knobs” by which to tune the simulation to measured data. Properties of the many materials alone provide hundreds more “knobs.” This means building model calibration is a severely under-determined problem with thousands of inputs that can be varied to match potentially only 12 data points, leading to questions as to the physical realism of the final tuned model compared with the very large number of valid models. As a starting point, this paper will focus on the most important task of getting the thermal characteristics of the envelope right by tuning envelope properties to match electricity consumption throughout the building.

It should come as no surprise that, to reduce the cost of doing business, the idea of developing self-calibrating energy models has been around

for decades, with initial attempts beginning around the early 1980s. Many of the motivations, history, different levels of thoroughness in calibration, and previous approaches (e.g., sensitivity analysis, reducing the number of simulations necessary, optimization methods) are expertly consolidated in ASHRAE report 1051-RP on the subject [18]. FEMP also has rigorous measurement and validation (M&V) guidelines [19].

### 2.6. Computational Complexity

The EnergyPlus whole-building energy simulation engine was consolidated from diverse sources around 1995, with functionality traceable back to DOE-2 and DOE's Building Loads Analysis and System Thermodynamics (BLAST) from the late 1970s [6]. Its original design goals were to provide a more consistent software structure for development and modification, to allow third-party programs and components to easily interface with the core system, and to fully integrate the loads, systems, and plants into the simulation [6]. The workflow for a building modeler using a system like EnergyPlus is to create a building's geometry using external software, layer it with detailed metrics encoding material properties, and add equipment currently in or expected to be in the building, including anticipated operational schedules.

A typical residential building model in EnergyPlus has approximately 3,000 input parameters that must be specified. The search space in such a problem is extremely large. Even if each parameter were a simple binary value (e.g., a categorical whose values were *yes* or *no*), the search space for a 3,000-parameter building would contain  $2^{3000}$  possibilities, which is unfathomably larger than the number of atoms believed to exist in the observable universe (about  $2^{270}$ ). This is a best-case scenario; the actual size of the search space

is effectively infinite because many of the parameters are continuous-valued.  
For this study, we relied on several building simulation experts to identify  
156 input parameters, along with minimums and maximums within which  
they will be varied. Even with this simplification, to solve this problem with  
brute force would take Titan more than  $10^{28}$  lifetimes of the known universe  
(13.8 billion years) to calculate all the necessary EnergyPlus simulations to  
find the best model for a single residential building. Many of the experiments  
in this paper discuss algorithms and methods to find near-optimal models in  
a tractable amount of time on commodity hardware.

### 2.7. Evolutionary Computation

A common approach to such search problems, and one that has been successfully applied in previous building model tuning work [10], is evolutionary computation (EC). EC [20, 21, 22, 23] is a stochastic search algorithm that attempts to mimic biological evolution by maintaining a set of candidate solutions (see Figure 1), referred to as a *population*. Each candidate solution is evaluated to determine its *fitness*, which is a problem-dependent measure of how well it solves the problem [24]. In essence, the candidate solutions act as samples of the search space, and their fitness values provide an approximation of the gradient. However, unlike strict gradient-based techniques, an EC provides some (typically small) opportunity for less fit candidates to influence the search process, which helps the EC to avoid local optima [25].

Since EC is a population-based search, it has a number of aspects that are inherently parallel. Most important for this work, each candidate in the population must be evaluated to determine its fitness, and the evaluation of one candidate is typically independent of the others. For expensive evaluation

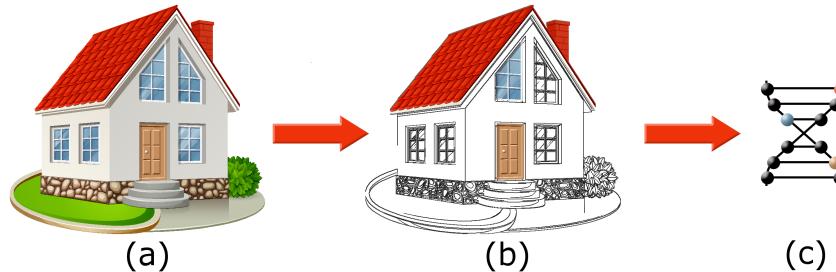


Figure 1: Candidate solutions represent the essential components of the problem that should be optimized. In the building-model-calibration case, the design of a candidate solution progresses as follows: (a) The base model is analyzed to determine optimization variables. (b) The optimization variables are isolated. (c) The optimization variables are encoded in some structure that can participate in the evolution (known as the *genotype*).

functions, the ability to determine the fitness values of an entire population simultaneously greatly increases the search efficiency. In addition to parallel evaluations, EC approaches called *island models* [26] exist that maintain multiple populations, each of which performs parallel evolution. In many cases, each island (i.e., isolated population) is allowed to explore a different region of the search space, possibly in different ways. With most island models, there are mechanisms in place that allow candidate solutions to *migrate* between islands (usually with low frequency), which allows information exchange between the populations.

In the context of this work, a candidate solution is a building with a chromosome represented by a list of building model parameters to be tuned. To evaluate a candidate solution, its corresponding EnergyPlus input data file (IDF) model is constructed and passed to the EnergyPlus simulation engine, producing an entire set of output measures (e.g., heating/cooling

loads). These output measures are then compared with actual measured data from the building. The resulting measure of accuracy is used as the fitness value for the candidate solution.

The interaction among candidate solutions in the population drives the evolutionary process, and this interaction is defined in terms of *evolutionary operators*. These operators determine how current solutions are recombined or modified to produce new solutions, as well as which solutions get to contribute information to the next generation. Figure 2 illustrates the basic concept of an evolutionary operator (recombination, in this case). The particular evolutionary operators used in this work were heuristic crossover, Gaussian mutation, tournament selection, and generational replacement. Each of these operators is standard in the EC literature, and all were described previously in [10].

### 3. Methodology

Previous work [10] focused on tuning a building to monthly utility data. The work presented here takes a more ambitious approach by attempting to optimize the match between a model building and actual hourly electricity usage data. The reference building used in this work is house number 1 in the Wolf Creek subdivision (WC1), an ORNL ZEBRAAlliance experimental energy-efficient home. This home has a plethora of energy-efficiency technologies: (1) standing seam metal roof with infrared reflective pigments to boost solar reflectance, (2) ENERGY STAR appliances, (3) triple-pane, low-emittance argon-filled windows, (4) compact fluorescent lighting, (5) a horizontal ground loop installation that leverages foundation and utility ex-

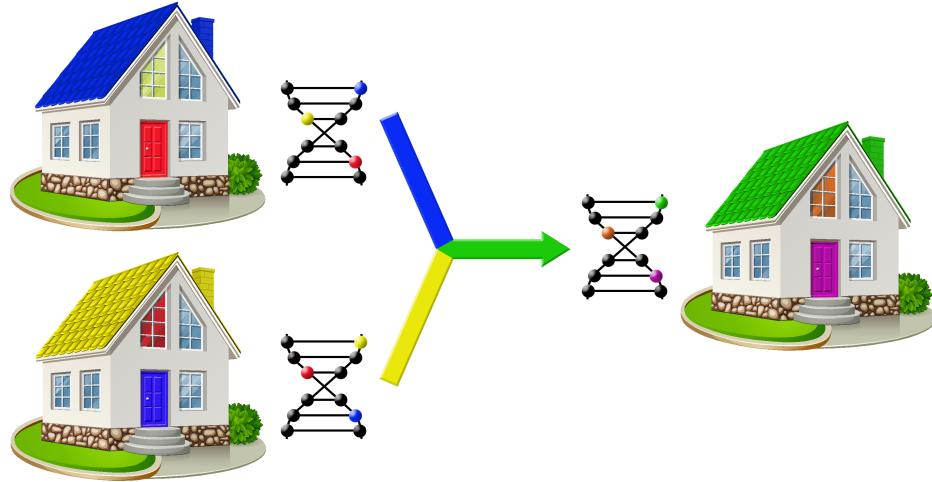


Figure 2: Evolution progresses as candidate solutions undergo the effects of evolutionary operators, such as recombination. In this illustration, the two genotypes (blue-yellow-red and yellow-red-blue) are recombined to produce an offspring candidate solution (green-orange-purple). The particular recombination illustrated here through color-mixing is similar to an approach in the EC literature known as *blend crossover*.

cavations, (6) high-efficiency water-to-air heat pump for space conditioning, (7) high-efficiency water-to-water heat pump for hot water heating, (8) an energy recovery ventilator to transfer heat and moisture between fresh incoming and outgoing air, and (9) structurally insulated panel (SIP) walls filled with expanded polystyrene insulation. For more information, the interested reader is referred to [11, 12]. The home has been fitted with more than 300 sensors that collect data at 15-minute timesteps using standard, wired-sensor data acquisition systems.

### 3.1. Base Models

The data, models, and accuracy metrics described in the remainder of this subsection are an extension of those previously reported in [10]. In the following experiments, two different model buildings are used. The first model, referred to here as the *refined* model, was last modified on March 29, 2012. This model matches whole-building annual electric consumption exactly, but it has a sum of absolute errors (SAE) of 1,276.34 kWh for monthly and 6,242.04 kWh for hourly electrical data over the entire year. An earlier version of that same model was completed on July 28, 2010, and it is referred to here as the *primitive* model. It has an SAE of 1,623.36 kWh for monthly and 8,113.69 kWh for hourly electrical data when summed for the entire year. These two baseline models are separated by approximately four person-months of effort (consisting of two months of laboratory material testing and two months manually tuning the input file) over the course of nearly two calendar years, with the refined model being the recipient of that effort.

1  
2  
3  
4  
5  
6  
7  
8  
9     3.2. Tuning Parameters  
10

11     Only a subset of the real-value parameters of the models, as specified by  
12 domain experts, was used as a part of the tuning process; the phrase “tuning  
13 parameters” will be used to refer to these variables. Although all 156 tunable  
14 parameters provided through the project website<sup>2</sup> are too extensive to list  
15 here, a majority of the parameters were for building material properties such  
16 as thickness, conductivity, density, specific heat, thermal absorptance, solar  
17 absorptance, and visible absorptance for materials such as gypsum board,  
18 stone, concrete foundation wall, fiberglass insulation, metal roofing, ply-  
19 wood, insulation, gravel, oriented strand board (OSB), and cladding, as well  
20 as U-factor, solar heat gain coefficient, and visible transmittance for window  
21 glazing systems. Other parameters include fraction of latent and radiant for  
22 equipment, radiant and visible for lighting, flow coefficients for HVAC, heat-  
23 ing and cooling air supply temperatures, building orientation, infiltration,  
24 and several others. Note also that while these are individual line changes in  
25 an EnergyPlus IDF input file, several instances of each material, equipment,  
26 and so on may be used throughout a building. The Autotune system cur-  
27 rently scales to tune any set of numerical parameters and can be customized  
28 for tuning only selected parameters in customizable ways. This allows cali-  
29 bration according to the needs of a particular use case or comparison with  
30 manual tuning efforts.

49  
50     

---

51         <sup>2</sup>[http://bit.ly/autotune\\_parameters](http://bit.ly/autotune_parameters)  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11     3.3. *Measured Data*

12     In 2010, the average US residential building consumed enough energy to  
13     cost the homeowner \$2,201 [1]. An average of 44.7% of the energy went to  
14     space heating and 9.2% toward space cooling [27], totaling 53.9% for space  
15     conditioning. However, primarily because of differing costs for various fuel  
16     types [28], 28.9% of the cost was for heating and 14.0% for cooling, yielding  
17     42.9% of the cost and amounting to \$944/year for space conditioning. How-  
18     ever, the energy-efficient HVAC in WC1 actually cost \$472.62 for January  
19     1–November 30, 2010. This 50% cost reduction for an energy-efficient home  
20     may serve as a point of reference for the tuning results presented throughout  
21     the residential study.

22     For the all-electric WC1, the actual energy usage data for all HVAC  
23     equipment was reliably collected from January 1 through November 28, 2010,  
24     at which point a new set of test HVAC equipment was installed. Therefore,  
25     in all experiments reported, the “yearly” electrical usage (and, likewise, the  
26     “full” schedule) will always refer to the electrical usage from January 1 to  
27     November 28. The electrical usage in this work was calculated as the sum of  
28     all of the heating and cooling ideal loads for every time period (in kilowatt-  
29     hours).

30  
31     3.4. *Tuning Accuracy*

32     In the following experiments, the primary metric used for measuring tun-  
33     ing accuracy is the SAE<sup>3</sup>. The SAE was calculated according to Eq. 1, where

34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54     <sup>3</sup>The SAE was chosen as the primary metric, rather than RMSE, for example, because  
55     of its ease of interpretation in terms of dollars difference.  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

1  
2  
3  
4  
5  
6  
7  
8  
9      $M_i$  is the heating+cooling load of the model and  $A_i$  is the heating+cooling  
10 load of the actual ZEBRAlliance WC1 building. This equation contains only  
11  $n = 8016$  hours because actual data was not collected for December; this  
12 value is indicated as  $SAE_{hour}$ . One of the acceleration methods used is to  
13 tune on an abbreviated schedule of only four days (January 1, April 1, Au-  
14 gust 1, and November 1) before transitioning to a full year; use of this variant  
15 of the metric will be indicated as four-day SAE.

22  
23  
24                       $SAE = \sum_{i=1}^n |M_i - A_i|$                       (1)

25  
26  
27     One practical consideration in moving from monthly to hourly data is  
28 that failures and sensor drift are more prevalent and easily detectable in  
29 hourly data. For this dataset, it was found that sensors failed to make a  
30 measurement in approximately 2% of all measurements. In [10], in dealing  
31 with error calculations, the failures were treated as 0 values in the summa-  
32 tion because it was believed they would not significantly impact monthly  
33 usage. For hourly data, however, they make a fourfold larger difference, and  
34 it was decided to ignore any hour that contained at least one sensor failure  
35 (approximately 8% of the data).

36  
37  
38     3.5. *Evolutionary Computation Parameters*

39  
40     In the following experiments, we always use 1,024 fitness evaluations (En-  
41 ergyPlus simulations compared with actual data) to allow comparison across  
42 experiments. These evaluations are organized into 16 individuals (build-  
43 ing models) evolved over 64 generations, tournament selection with tourna-  
44 ment size 4, generational replacement with weak elitism (one elite), heuristic  
45

1  
2  
3  
4  
5  
6  
7  
8  
9 crossover, and Gaussian mutation with a usage rate of 1.0 and a mutation  
10 rate of 10% of the allowable range of each variable, unless otherwise specified.  
11 Because of the stochastic nature of this approach, results are shown for eight  
12 independent trials.  
13  
14

15 **4. Experiments**  
16  
17

18 *4.1. Experiment 1—Determining a Fitness Surrogate*  
19  
20

21 An annual EnergyPlus simulation for the WC1 reference building at 1  
22 hour time intervals (over 8,000 time periods) requires 8 minutes of runtime  
23 and cannot currently be efficiently parallelized. This type of fine-grained  
24 schedule leads to a much more exact, but computationally expensive, sim-  
25 ulation. A 4 day simulation has only 96 time periods of output data but  
26 can run in seconds. It was also previously established [10] that tuning on  
27 an abbreviated schedule of only 4 days (January 1, April 1, August 1, and  
28 November 1) was a defensible acceleration method for monthly data, and  
29 this experiment validates the same behavior for hourly data. We do this by  
30 establishing a correlation between the error rate of the 4 day and full sched-  
31 ules, which requires a particular kind of sampling spanning from high to low  
32 error. An EC was used to minimize the error in the abbreviated schedule,  
33 and the hourly error rates in the 4 day and full schedules of each created  
34 individual were measured and stored to determine their correlation. To en-  
35 sure against statistical anomalies, the EC was run four different times with  
36 different initial populations each time.  
37  
38

39 The entire set of 1,024 individuals for each of the four trials is plotted in  
40 Figure 3. This plot compares the hourly sum of the absolute errors (SAE)  
41  
42

1  
2  
3  
4  
5  
6  
7  
8  
9 between the candidate solutions and the actual energy usage, for both the 4  
10 day schedule and the full schedule, for each of the four independent trials. It  
11 is clear from the graph that there is a strong linear correlation between the  
12 electricity usage with the 4 day schedule and that with the full schedule. More  
13 precisely, this correlation was 0.9697, 0.9704, 0.9472, and 0.9382 for each of  
14 the four trials, respectively. This can be seen in the first row of Table 1,  
15 which shows how the correlation changes as the data focus more on the  
16 later generated samples (which would typically have lower SAE). So, as with  
17 previous results, the 4 day abbreviated schedule appears to be an acceptable  
18 surrogate for the full schedule. An open, relevant, and interesting question  
19 is what is the smallest subset of days that provides the best correlation to  
20 full-schedule error.

31  
32  
33 *4.2. Experiment 2—Tuning Using the Abbreviated Schedule*

34  
35 The abbreviated schedule is desirable since it is significantly less expensive  
36 to compute, has been shown to be an effective surrogate for full-schedule  
37 error, and provides a baseline against which other results can be compared  
38 and contextualized. In Experiment 2, we tune the refined and primitive  
39 models on hourly electrical usage with the abbreviated schedule. We then  
40 use the 4 day tuned models to calculate data for the entire year to get a  
41 full-schedule SAE that is directly comparable to other experiments.

42  
43  
44 Table 2 presents the final population statistics for full-schedule SAE on  
45 each model and trial. The actual full-schedule data were compared with each  
46 of the 16 candidates in the final population for each trial, and the average  
47 and minimum SAE values from this comparison are reported in this table.  
48 For the refined model, the average minimum yearly SAE achieved by the EC

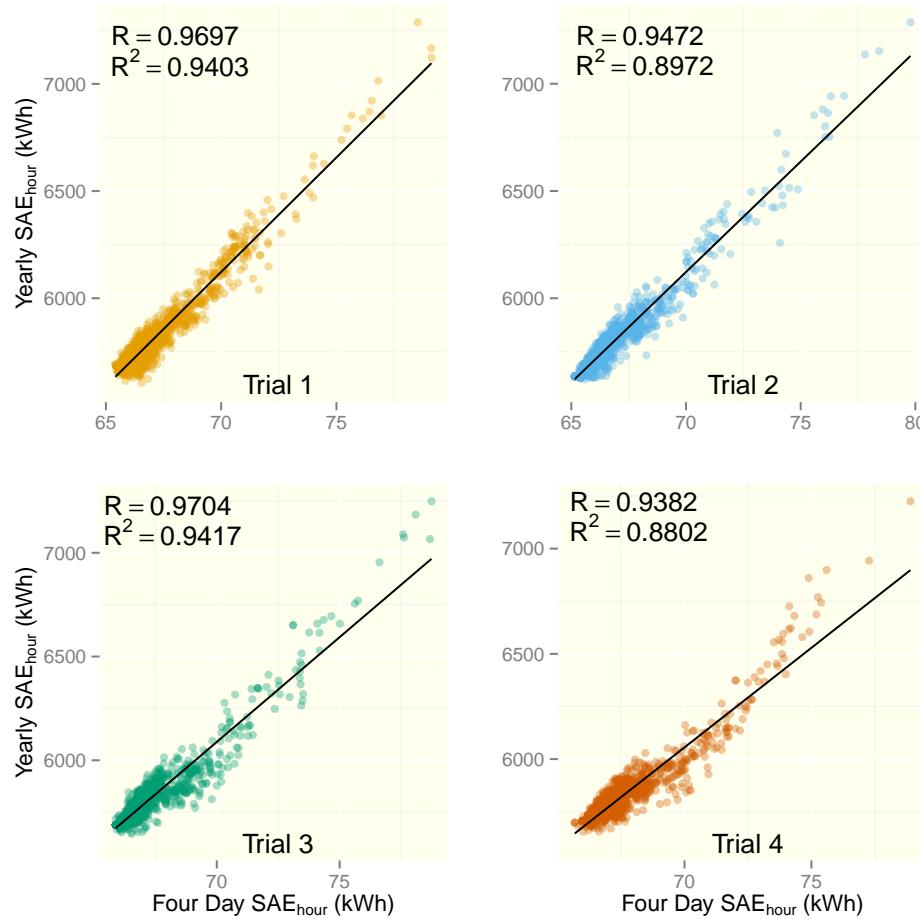


Figure 3: Each of the four independent trials shows a very high correlation between 4 day and full-year-schedule SAE of hourly electrical usage. This means that an abbreviated surrogate (4 -day schedule) is a plausible method of expediting the tuning.

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25

Table 1: The sampling of the SAE to determine correlation was done using an EC with the objective of minimizing the SAE. This produced initial samples with relatively high SAE values and subsequent samples with generally lower SAE as more data were used. The correlation using only the first  $N\%$  of the samples produced shows that even the initial samples are strongly correlated most of the time (with the exception of Trial 3).

26

Data Used	Trial 1	Trial 2	Trial 3	Trial 4
100%	0.9603	0.9422	0.9015	0.9555
90%	0.8668	0.8438	0.6545	0.8788
80%	0.8683	0.8573	0.6444	0.8922
70%	0.8726	0.8561	0.6231	0.8944
60%	0.8671	0.8489	0.6078	0.8952
50%	0.8703	0.8489	0.5557	0.8971
40%	0.8774	0.8417	0.5663	0.9020
30%	0.8794	0.8208	0.6015	0.8973
20%	0.8841	0.8231	0.5988	0.8860
10%	0.8796	0.8235	0.3985	0.8714

was 5,660 kWh. This corresponds to a reduction of 582 kWh in SAE—an almost 10% reduction from the model, which had an SAE of 6,242 kWh. For the primitive model, the average minimum yearly SAE was 7,453 kWh, which is a reduction of 660 kWh or an 8% reduction in SAE. This is also a 35% reduction in the error between the primitive and refined models (8,114 kWh for the primitive and 6,242 kWh for the refined).

Table 2: In Experiment 2, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the abbreviated schedule. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7539	7578	7683	7638	7495	7467	7503	7587
Refined	Avg	5720	5708	5792	5762	5750	5783	5806	5698
Primitive	Min	7393	7507	7595	7545	7393	7360	7356	7477
Refined	Min	5620	5625	5709	5687	5640	5671	5711	5616

### 4.3. Experiment 3—Tuning Using the Full Schedule

As a means of comparison, this experiment focuses on using the full schedule for tuning the models. This approach is more computationally expensive than using the abbreviated schedule. The only difference between this experiment and Experiment 2 is the use of the full schedule for the EnergyPlus simulation for calculating candidate fitness. The expectation is that this

would generate tuned models with lower SAE values when compared against  
actual data.

Table 3 displays the final population statistics for the full-schedule SAE  
on each model and trial. As in Table 2, these statistics are the population  
averages and minimums. For the refined model, the average minimum SAE  
across eight trials was 5,539 kWh. Compare this error with that of the  
baseline file (6,242 kWh) and of the abbreviated schedule EC (5,660 kWh).  
Clearly, using the full schedule throughout the evolution provides a lower  
error (by 121 kWh) than using the abbreviated schedule. However, this  
additional 2% reduction in error beyond that of the abbreviated schedule  
requires more than five times the computational resources. The abbreviated-  
schedule EC was able to complete a single trial in about 1.5 hours on an  
eight-core machine, whereas the full schedule EC required 8 hours on the  
same machine. For the primitive model, the average minimum SAE across  
eight trials was 7,162 kWh. Once again, this error must be compared with  
the baseline file (8,114 kWh) and with the abbreviated-schedule EC (7,453  
kWh). The result shows a reduction of 291 kWh in SAE from the abbreviated  
schedule, which corresponds to a 4% reduction for a similar fivefold increase  
in computational time.

#### 4.4. Experiment 4—Tuning Using Both Schedules Serially

As shown in Experiment 3, the full schedule produces better tuning re-  
sults but requires more computational resources. If additional resources are  
allotted, it may be best to employ those after the tuning process has pro-  
duced a set of faithful models. The EC can be started from a known set  
of good solutions (which we call *seeds*), and that allows us to tune in two

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

Table 3: In Experiment 3, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the full schedule. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7275	7245	7250	7259	7338	7202	7274	7272
Refined	Avg	5582	5641	5620	5594	5636	5623	5623	5596
Primitive	Min	7175	7138	7150	7122	7253	7114	7162	7179
Refined	Min	5518	5568	5527	5510	5560	5550	5558	5523

parts—first with the abbreviated schedule to produce a set of seeds and then with the full schedule to produce the final tuning.

In this experiment, the first 768 E+ simulations were done using the 4 day schedule. At the end of those simulations, the best half (8) of the final population’s candidate solutions were inserted as seeds into a new initial population of 16 candidates. (The other 8 candidates were randomly generated.) This population was then evolved for the remaining 256 evaluations on the yearly schedule.

Table 4 shows the final population minimums and averages for each trial for both the refined and primitive models. The average minimum SAE for the refined model was 5,581 kWh, and for the primitive model was 7,343 kWh. As expected, the performance is better than that using the abbreviated schedule alone (5,660 kWh and 7,453 kWh for refined and primitive models,

respectively), but it is not as good as using the full schedule throughout the evolution (5,539 kWh and 7,162 kWh). The reduction in SAE was 79 kWh for the refined model and 110 kWh for the primitive model, which corresponds to a movement of 66% and 38% toward the full schedule performance for each model. This performance comes at a cost of approximately double the computational expense of the abbreviated schedule only (3 hours instead of 1.5 hours), which is less than the fivefold increase of the full schedule (over 8 hours).

Table 4: In Experiment 4, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the abbreviated schedule for the first 75% of the evaluations and the full schedule for the remaining 25%. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7404	7463	7469	7467	7393	7381	7434	7403
Refined	Avg	5626	5670	5688	5614	5695	5667	5642	5646
Primitive	Min	7306	7382	7402	7378	7328	7283	7362	7306
Refined	Min	5564	5586	5608	5545	5596	5618	5566	5563

#### 4.5. Experiment 5—Tuning Using Both Schedules in Parallel

Rather than executing the different evolutionary approaches in serial fashion, it is possible to perform them in parallel using an island-model EC [26]. In this experiment, two different populations of 16 candidate solutions are

1  
2  
3  
4  
5  
6  
7  
8  
9 maintained. The first (“fast”) island uses the abbreviated schedule for sim-  
10 ulation, and the second (“slow”) uses the full-year schedule. At the end of  
11 each generation (i.e., each 16 simulations), a one-element migration queue is  
12 queried for any existing candidate. If one is found, it replaces the worst can-  
13 didate in the population. In either case, the best candidate in the population  
14 is inserted into the queue. Since both populations perform this migration,  
15 this enables solutions to transfer between the populations (or “islands”).  
16 Note that that all immigrant solutions are re-evaluated to determine the fit-  
17 ness under the new simulation approach before they are inserted into the  
18 population. This island-model approach is illustrated in Figure 4.  
19

20 Table 5 shows the final population statistics for the candidates from the  
21 abbreviated-schedule islands for both primitive and refined models. The av-  
22 erage minimum SAE was 5,597 kWh for the refined model and 7,270 kWh  
23 for the primitive model. This is very near the performance obtained from the  
24 serial evolution. However, the parallel approach has a number of advantages.  
25 First, and perhaps most important, it makes no assumptions about the num-  
26 ber of simulations that will be carried out, so users are free to stop the tuning  
27 process at any time. Since the full simulation is always being run on one of  
28 the islands, this information is always incorporated in the final population.  
29 Second, the parallel approach appears to provide slightly better performance  
30 in terms of the SAE-to-full-year-simulation ratio. This is because the actual  
31 performance between the two is nearly identical, but the parallel version’s  
32 full island was unable to make use of particularly good solutions from the  
33 abbreviated island until after about the first 25% of simulations had been  
34 run. So, in some sense, it achieved similar error performance using only 75%  
35

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

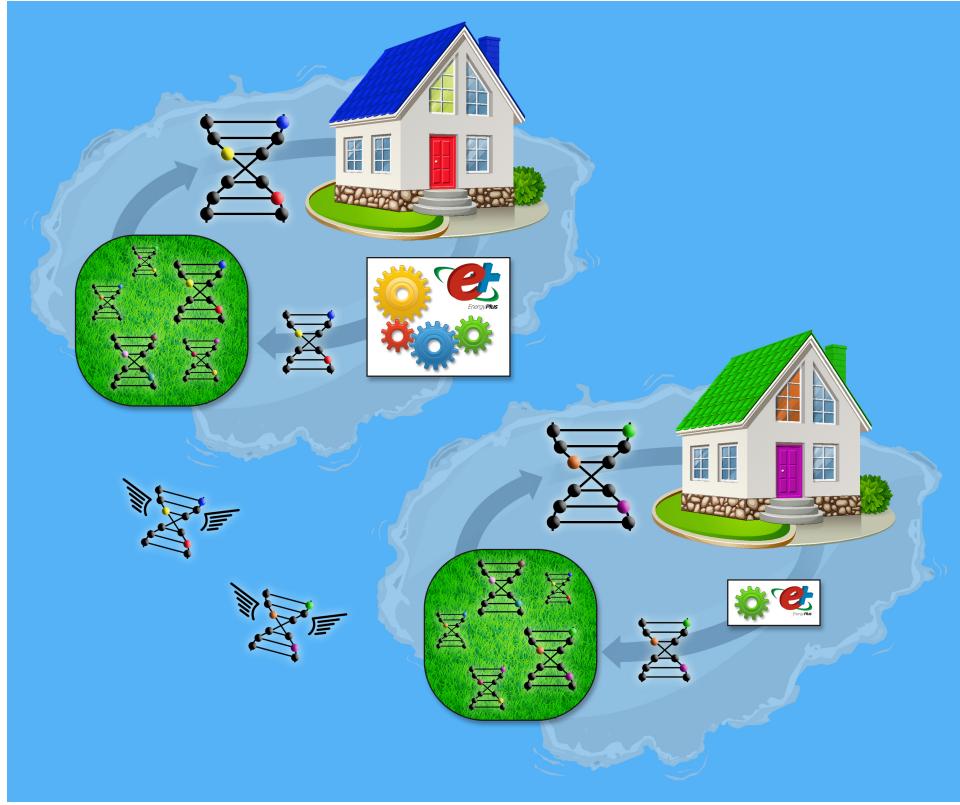


Figure 4: When performing the tuning of both schedules in parallel, two different populations evolve independently. The evolution on each island proceeds by choosing a candidate solution, generating the EnergyPlus model based on that genotype, passing that model through the EnergyPlus simulation, determining the candidate’s fitness (represented here by candidate solution size), and finally re-integrating the candidate solution into the population. However, the calculation of the fitness is handled differently on each island. One island (the upper-left) uses the full year schedule (represented by the large EnergyPlus “black box”), while the other island (the lower-right) uses the abbreviated schedule. Periodically during the evolution, candidate solutions are allowed to migrate between the islands (represented here as winged candidate solutions *en route*). When these candidate solutions arrive at their destination islands, they must be re-evaluated in order to determine their appropriate size (since their current sizes rely on the measure as defined on their home island).

of the simulations. This 25% is an estimate based on the typical convergence of the abbreviated island as shown in Figure 5. This figure plots the convergence of the error of the best refined-model candidate through each generation for each trial of Experiment 4 (up through generation 48, after which the full-schedule information was used). Generation 12 represents the 25% mark (shown as a dotted line in the figure) of the full 1,024 evaluations.

Table 5: In Experiment 5, an island-model evolutionary computation was used to optimize the primitive and refined models eight independent times (thus 16 different tunings) using a *slow* and a *fast* island. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7340	7330	7435	7470	7391	7349	7390	7356
Refined	Avg	5724	5675	5670	5660	5648	5715	5669	5625
Primitive	Min	7218	7201	7360	7373	7281	7211	7293	7226
Refined	Min	5639	5597	5582	5578	5580	5633	5591	5573

## 5. Conclusions and Future Work

The results of the previous experiments are summarized in Table 6 and Figure 6. This table shows the baseline and various approaches for tuning for each model, focusing on the SAE and RMSE between the model output and the actual data. The percentages in parentheses are the percent reduction in error from the baseline.

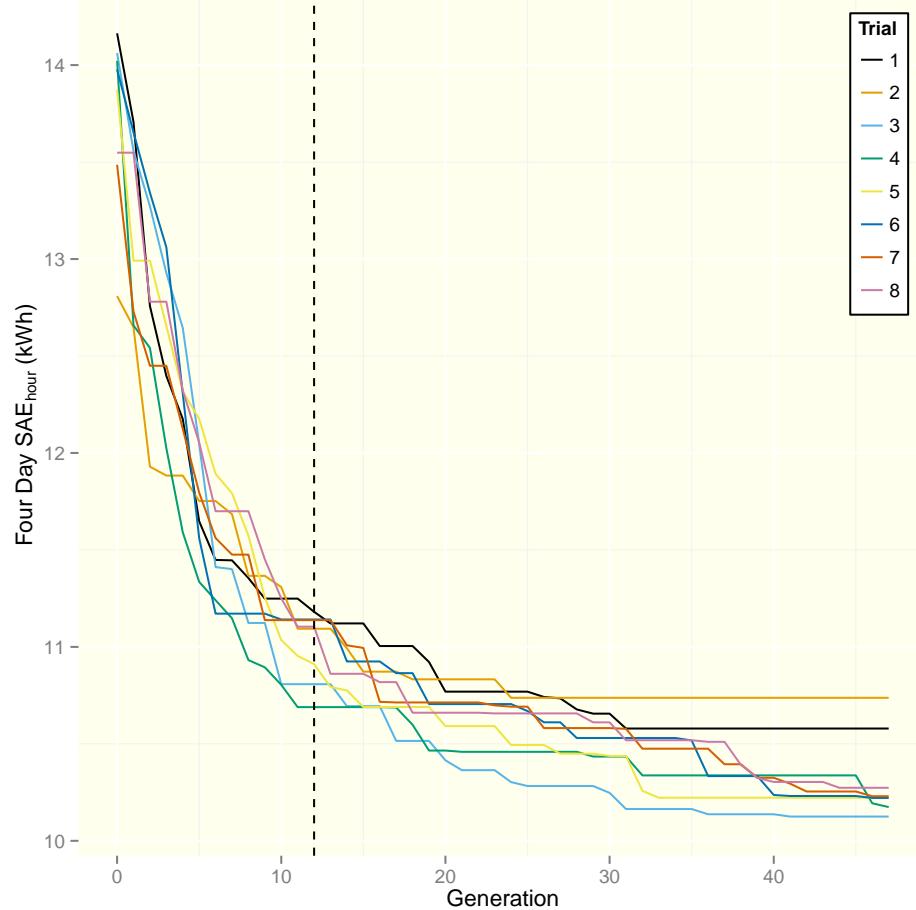


Figure 5: This figure shows the convergence of the best individual in each trial of Experiment 4 when the abbreviated schedule is used (i.e., the first 75% of the evaluations). The dotted line marks generation 12, which is 25% of the evaluations allotted to the abbreviated schedule. Before this point, individuals produced by the abbreviated schedule tuning are not particularly useful as seeds because of their relatively high SAE. After this point, they become much more useful.

1  
2  
3  
4  
5  
6  
7  
8  
9 These experiments have shown that it is possible to use an evolutionary  
10 search to reduce the RMSE by over 7% in a comparison of hourly electrical  
11 usage. They have also shown that a two-island approach is both a feasible and  
12 an effective way of combining the abbreviated- and full-schedule approaches.  
13  
14

15 Experiment 1 verified that a 4 day schedule could serve as a viable surro-  
16 gate for the full-year schedule in terms of hourly electrical usage, producing  
17 a correlation coefficient of 0.96 between abbreviated and full schedules. Ex-  
18 periment 2 showed that an abbreviated-only approach was able to reduce the  
19 RMSE by 6% and 7% below the baseline for the refined and primitive models,  
20 respectively. Experiment 3 focused on using the full schedule only, and that  
21 approach reduced the RMSE by 7% (refined) and 10% (primitive). Experi-  
22 ment 4 combined the two approaches in serial fashion, using the abbreviated  
23 schedule for the first 768 simulations and then using the full schedule for  
24 the remaining 256. This approach was better than the abbreviated schedule  
25 alone, but it performed only slightly worse than the full schedule alone. Fi-  
26 nally, Experiment 5 combined the two approaches in parallel using an island  
27 model, where one island used the abbreviated schedule and the other used  
28 the full schedule. Solutions were able to migrate between the islands, and  
29 this approach was competitive with the full schedule approach, despite the  
30 fact that it used only 25% of the full-schedule simulations.  
31  
32

33 In conclusion, this paper documents a series of experiments that elucidate  
34 several methods, advantages, and trade-offs in the use of EC to automatically  
35 calibrate a software model of an emulated-occupancy experimental residential  
36 building to hourly whole-building electrical data. Work is ongoing to extend  
37 these methods for calibration from residential buildings, where fewer things  
38  
39

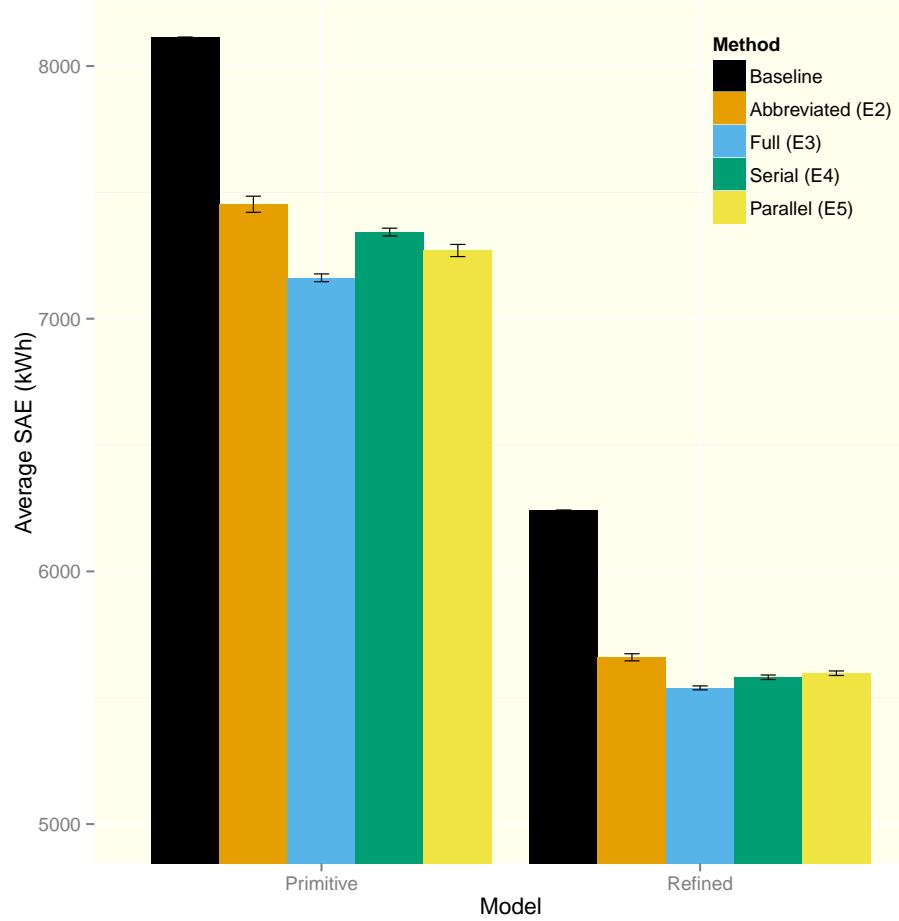


Figure 6: This figure shows the results of Experiments 2–5 compared with the baseline (see Table 6). The average minimum SAE value found across all eight trials is used to represent each method, and each method is accompanied by standard error bars (except for the baseline). The y-axis minimum value was increased to 5,000 to highlight the differences between the methods.

1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15

Table 6: The summary comparison of results from Experiments 2–5 show that significant reductions can be gained from tuning with an abbreviated schedule, either alone, serially, or in parallel with full-schedule tuning. Percentages in parentheses are reductions versus the baseline.

Approach	Refined		Primitive	
	SAE (kWh)	RMSE	SAE (kWh)	RMSE
Baseline	6242	1.206	8114	1.625
Abbreviated	5660 (9%)	1.129 (6%)	7453 (8%)	1.514 (7%)
Full	5539 (11%)	1.119 (7%)	7162 (12%)	1.458 (10%)
Serial	5581 (11%)	1.123 (7%)	7343 (9%)	1.497 (8%)
Parallel	5597 (10%)	1.121 (7%)	7270 (10%)	1.482 (9%)

30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

are certain, to larger commercial buildings, calibrating to sub-hourly data for multiple data channels and to quantitative tests for determining not only how closely the building matches measured data but also how accurately the tuned model matches the real building.

## 6. Acknowledgements

This work was funded by field work proposal CEBT105 under DOE Building Technology Activity Number BT0201000. We would like to thank Amir Roth for his support and review of this project. This research used resources of the Oak Ridge Leadership Computing Facility at ORNL, which is supported by the Office of Science of the DOE under Contract No. DE-AC05-00OR22725. Our work has been enabled and supported by data analysis and visualization experts at the RDAV (Remote Data Analysis and Visu-

alization) Center of the University of Tennessee-Knoxville (NSF grant no. ARRA-NSF-OCI-0906324 and NSF-OCI-1136246). ORNL is managed by UT-Battelle, LLC, for DOE under contract DE-AC05-00OR22725. This manuscript has been authored by UT-Battelle, LLC, under Contract Number DEAC05-00OR22725 with DOE. The United States Government retains and the publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paid-up, irrevocable, world-wide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes.

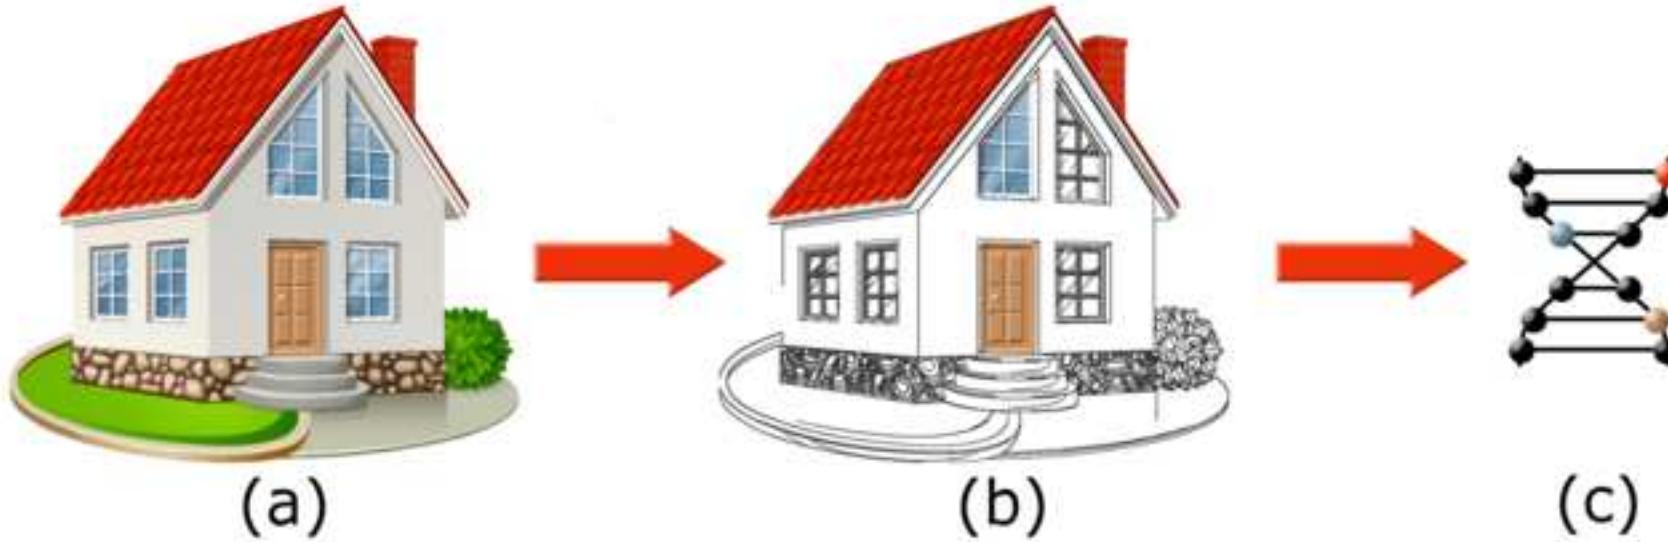
- [1] Department of Energy, Average annual energy expenditures per household, by year, in: Building Energy Data Book, 2012.
- [2] Earth Advantage Institute and Conservation Services Group, Energy performance score 2008 pilot: Findings and recommendations report.
- [3] D. Roberts, N. Merket, B. Polly, M. Heahey, S. Casey, J. Robertson, Assessment of the u.s. department of energy scoring tool, Tech. rep., National Renewable Energy Laboratory (2012).
- [4] ASHRAE, Guideline 14-2002: Measurement of Energy and Demand Savings, ASHRAE (2002).
- [5] Department of Energy Tools, Building energy software tools directory, available online at [http://apps1.eere.energy.gov/buildings/tools\\_directory/subjects\\_sub.cfm](http://apps1.eere.energy.gov/buildings/tools_directory/subjects_sub.cfm) (2012).
- [6] EnergyPlus, Getting Started with EnergyPlus: Basic Concepts Manual (October 2011).

- 1  
2  
3  
4  
5  
6  
7  
8  
9 [7] National Renewable Energy Laboratory, Openstudio: Commercial  
10 buildings research and software development, available online at <http://openstudio.nrel.gov> (2012).  
11  
12  
13  
14  
15 [8] ASHRAE, Handbook of Fundamentals (chapter 26): Heat, Air, and  
16 Moisture Control in Building Assemblies Material Properties, ASHRAE  
17 (2013).  
18  
19  
20  
21 [9] J. R. New, J. Sanyal, M. S. Bhandari, S. S. Shrestha, Autotune e+  
22 building energy models, in: Proceedings of the 5th National SimBuild  
23 of IBPSA-USA, International Building Performance Simulation Associa-  
24 tion (IBPSA), 2012.  
25  
26  
27  
28  
29  
30 [10] A. Garrett, J. New, T. Chandler, Evolutionary tuning of building models  
31 to monthly electrical consumption, in: ASHRAE Transactions, 2013.  
32  
33  
34  
35 [11] W. A. Miller, ZEBRAlliance: An alliance maximizing cost-effective en-  
36 ergy efficiency in buildings, Tech. rep., Oak Ridge National Laboratory  
37 (2012).  
38  
39 URL <http://www.zebralliance.org>  
40  
41  
42  
43  
44 [12] K. Biswas, A. Gehl, R. Jackson, P. Boudreaux, J. Christian, Comparison  
45 of two high-performance energy efficient homes: Annual performance  
46 report, Tech. rep., Oak Ridge National Laboratory (2012).  
47  
48  
49  
50 [13] J. Sanyal, J. R. New, Supercomputer assisted generation of machine  
51 learning agents for the calibration of building energy models, in: Pro-  
52 ceedings of Extreme Science and Engineering Discovery Environment  
53 (XSEDE 13), San Diego, CA, 2013.  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

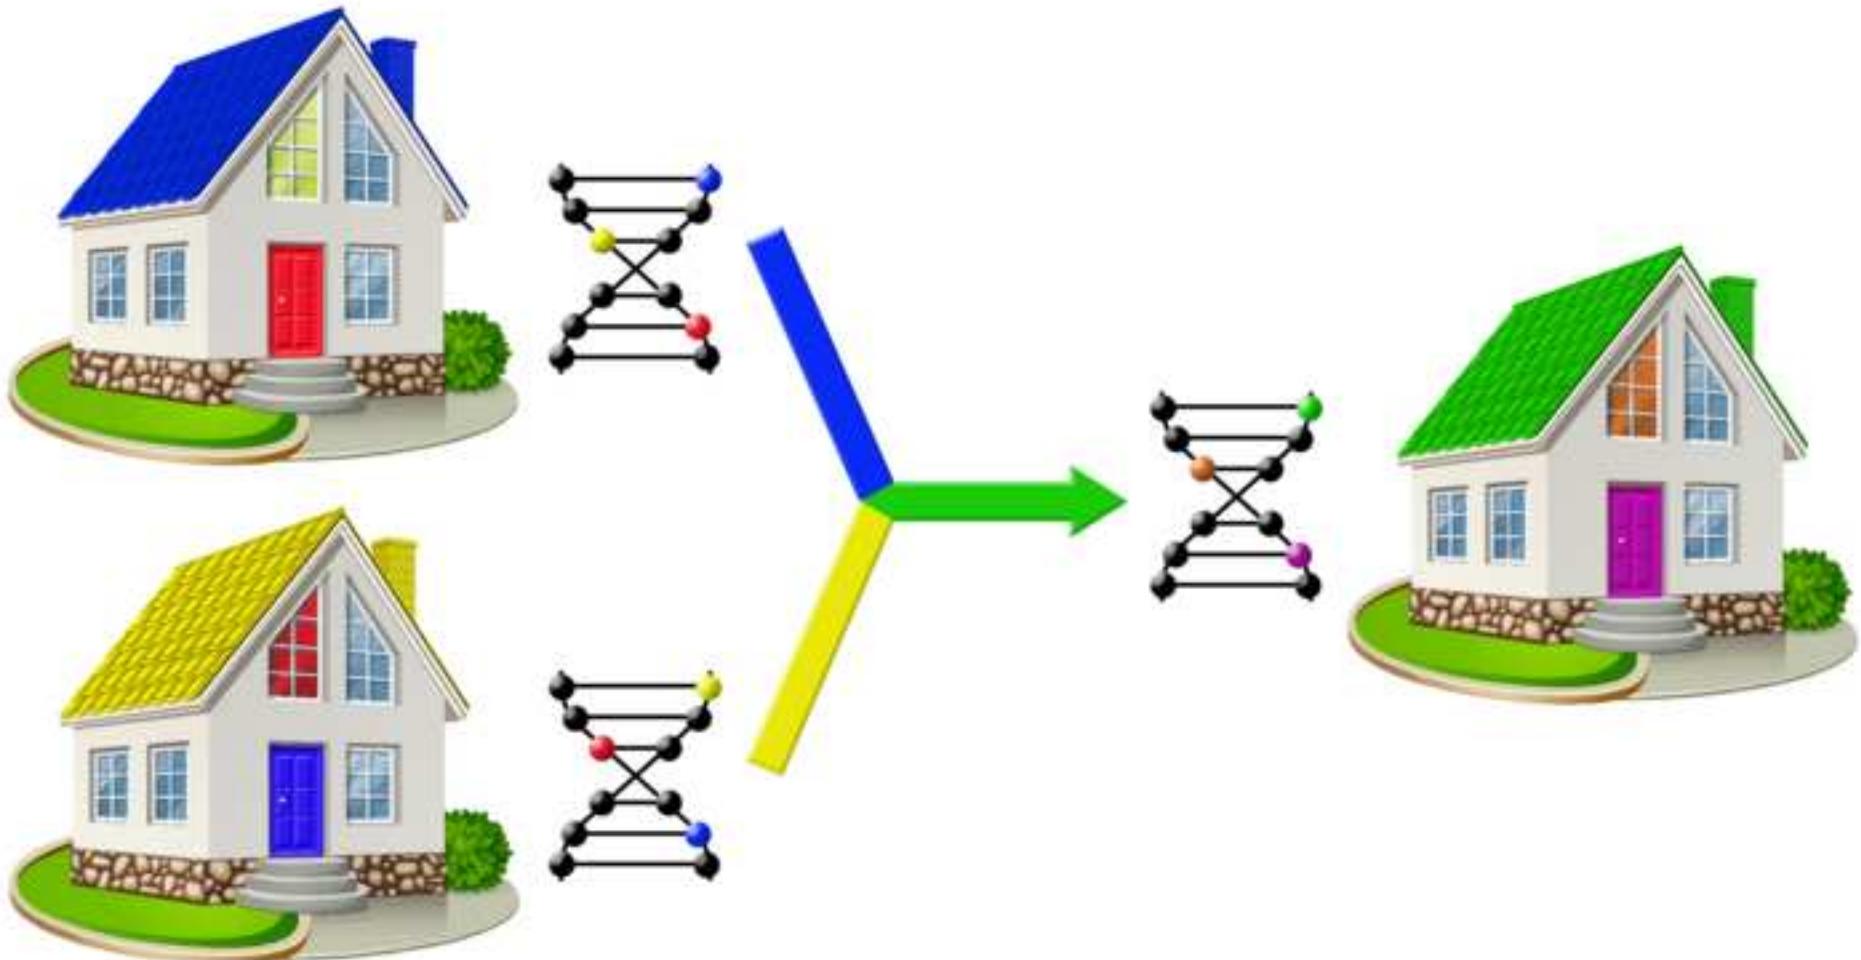
- 1  
2  
3  
4  
5  
6  
7  
8  
9 [14] J. Sanyal, J. R. New, Simulation and big data challenges in tuning building  
10 energy models, in: Proceedings of the IEEE Workshop on Modeling  
11 and Simulation of Cyber-Physical Energy Systems, Berkeley, CA, 2013.  
12  
13  
14  
15 [15] C. C. Castello, J. R. New, Autonomous correction of sensor data applied  
16 to building technologies utilizing statistical processing methods,  
17 in: Proceedings of the 2nd Energy Informatics Conference, 2012.  
18  
19  
20  
21 [16] R. E. Edwards, J. R. New, L. E. Parker, Constructing large scale energyplus  
22 surrogates from big data, Tech. rep., Oak Ridge National Laboratory  
23 (2013).  
24  
25  
26  
27 [17] S. Pigg, A rating tale, Home Energy Magazine.  
28  
29  
30 [18] T. A. Reddy, M. I., S. Jian, C. Panjapornporn, Procedures for reconciling  
31 computer-calculated results with measured energy data, Tech. rep.,  
32 ASHRAE (2006).  
33  
34  
35  
36 [19] Federal Energy Management Program, M&V Guidelines: Measurement  
37 and Verification for Federal Energy Projects Version 3.0, U.S. Department  
38 of Energy (apr 2008).  
39  
40  
41  
42 [20] K. A. DeJong, W. Spears, On the state of evolutionary computation,  
43 in: S. Forrest (Ed.), Proceedings of the Fifth International Conference  
44 on Genetic Algorithms, Morgan Kaufman, San Mateo, CA, 1993, pp.  
45 618–623.  
46  
47  
48  
49  
50  
51  
52 [21] W. M. Spears, K. A. DeJong, T. Bäck, D. B. Fogel, H. deGaris, An  
53 overview of evolutionary computation, in: Proceedings of the 1993 European  
54 Conference on Machine Learning, 1993.  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

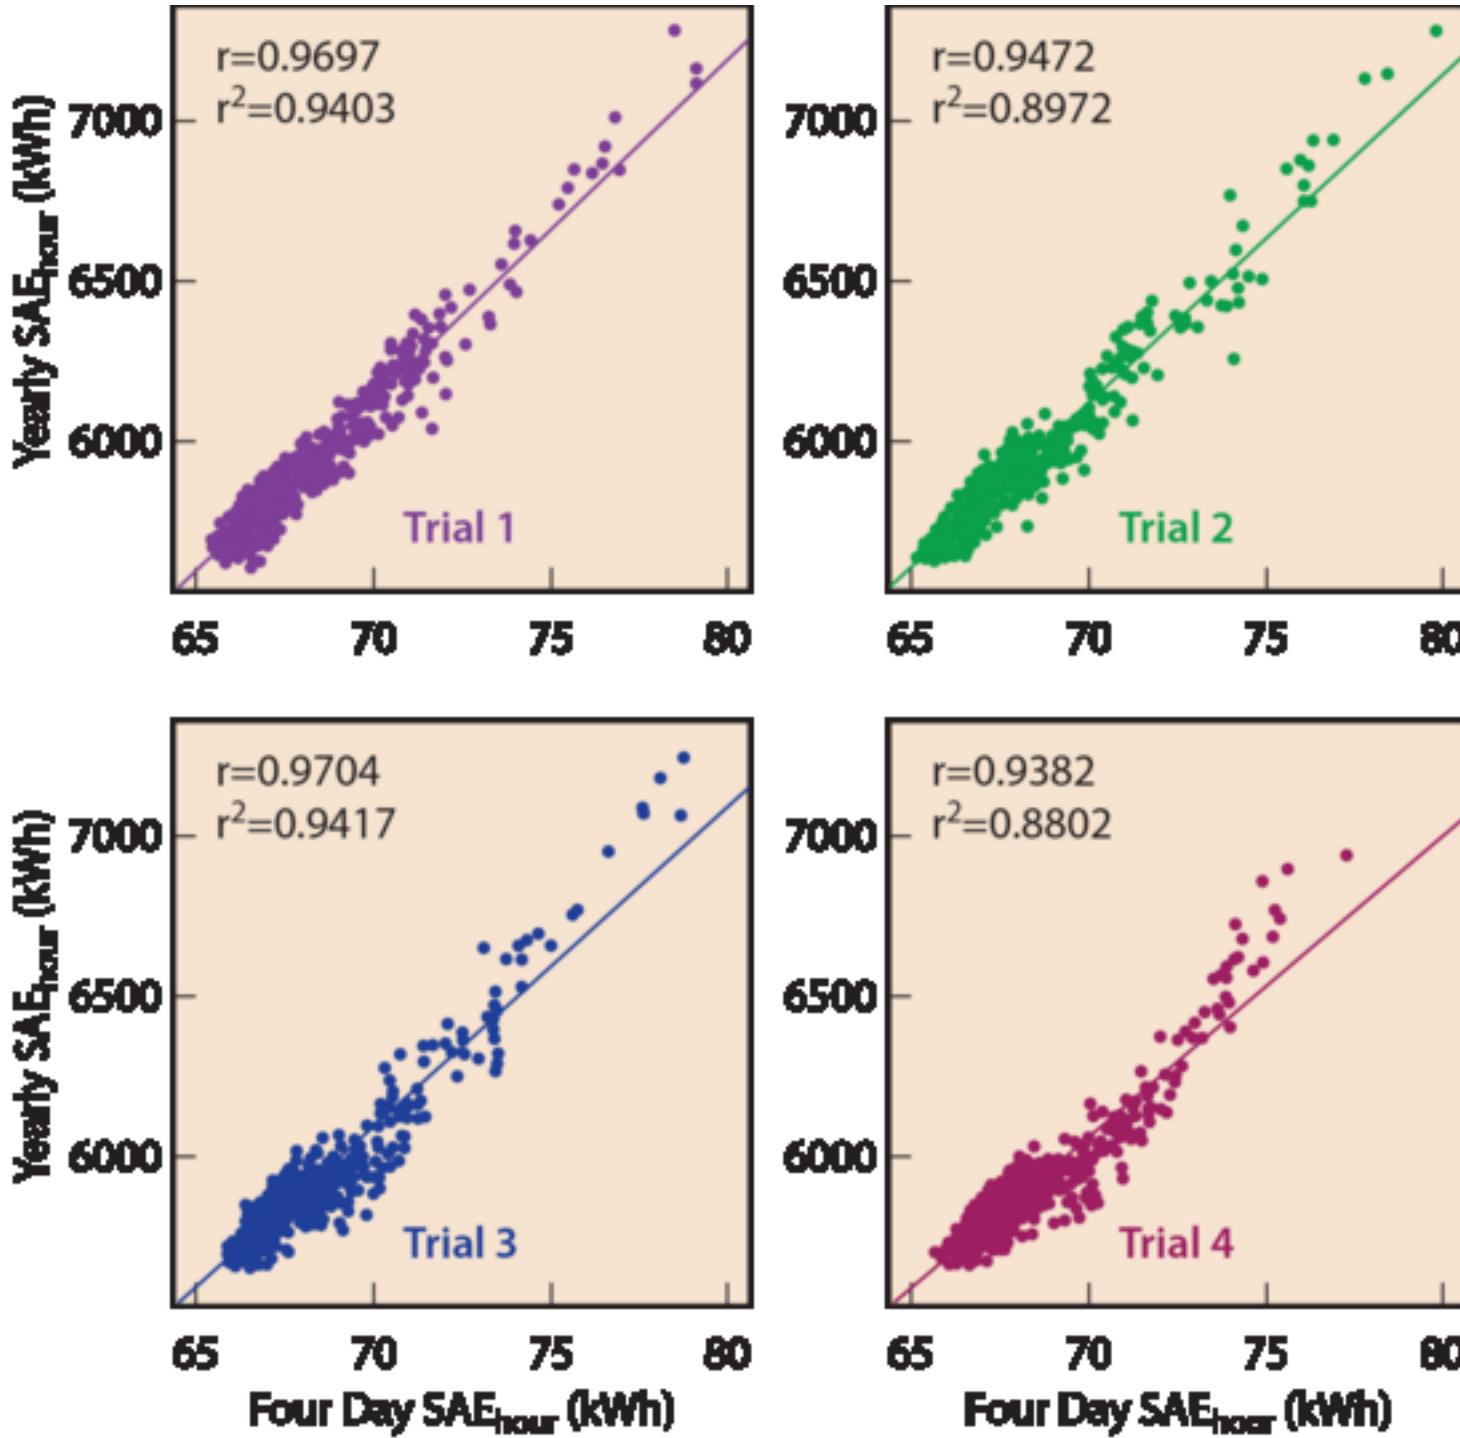
- 1  
2  
3  
4  
5  
6  
7  
8  
9 [22] D. B. Fogel, An introduction to simulated evolutionary optimization,  
10 IEEE Transactions on Neural Networks 5 (1) (1994) 3–14.  
11  
12  
13 [23] D. B. Fogel, What is evolutionary computation?, IEEE Spectrum 37 (2)  
14 (2000) 26–32.  
15  
16  
17  
18 [24] K. A. DeJong, Evolutionary Computation: A Unified Approach, MIT  
19 Press, 2006.  
20  
21  
22  
23 [25] Z. Michalewicz, D. B. Fogel, How to Solve It: Modern Heuristics,  
24 Springer, 2004.  
25  
26  
27  
28 [26] A. E. Eiben, J. E. Smith, Introduction to Evolutionary Computing,  
29 Springer, New York, 2007.  
30  
31  
32 [27] Department of Energy, Residential sector energy consumption. table  
33 2.1.5, in: Building Energy Data Book, 2012.  
34  
35  
36  
37 [28] Department of Energy, 2010 residential energy end-use expenditure  
38 splits, by fuel type, in: Building Energy Data Book, 2012.  
39  
40  
41  
42  
43  
44  
45  
46  
47  
48  
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65

How to turn a building into a chromosome  
[Click here to download high resolution image](#)



How to breed buildings (\*.psd available)  
[Click here to download high resolution image](#)

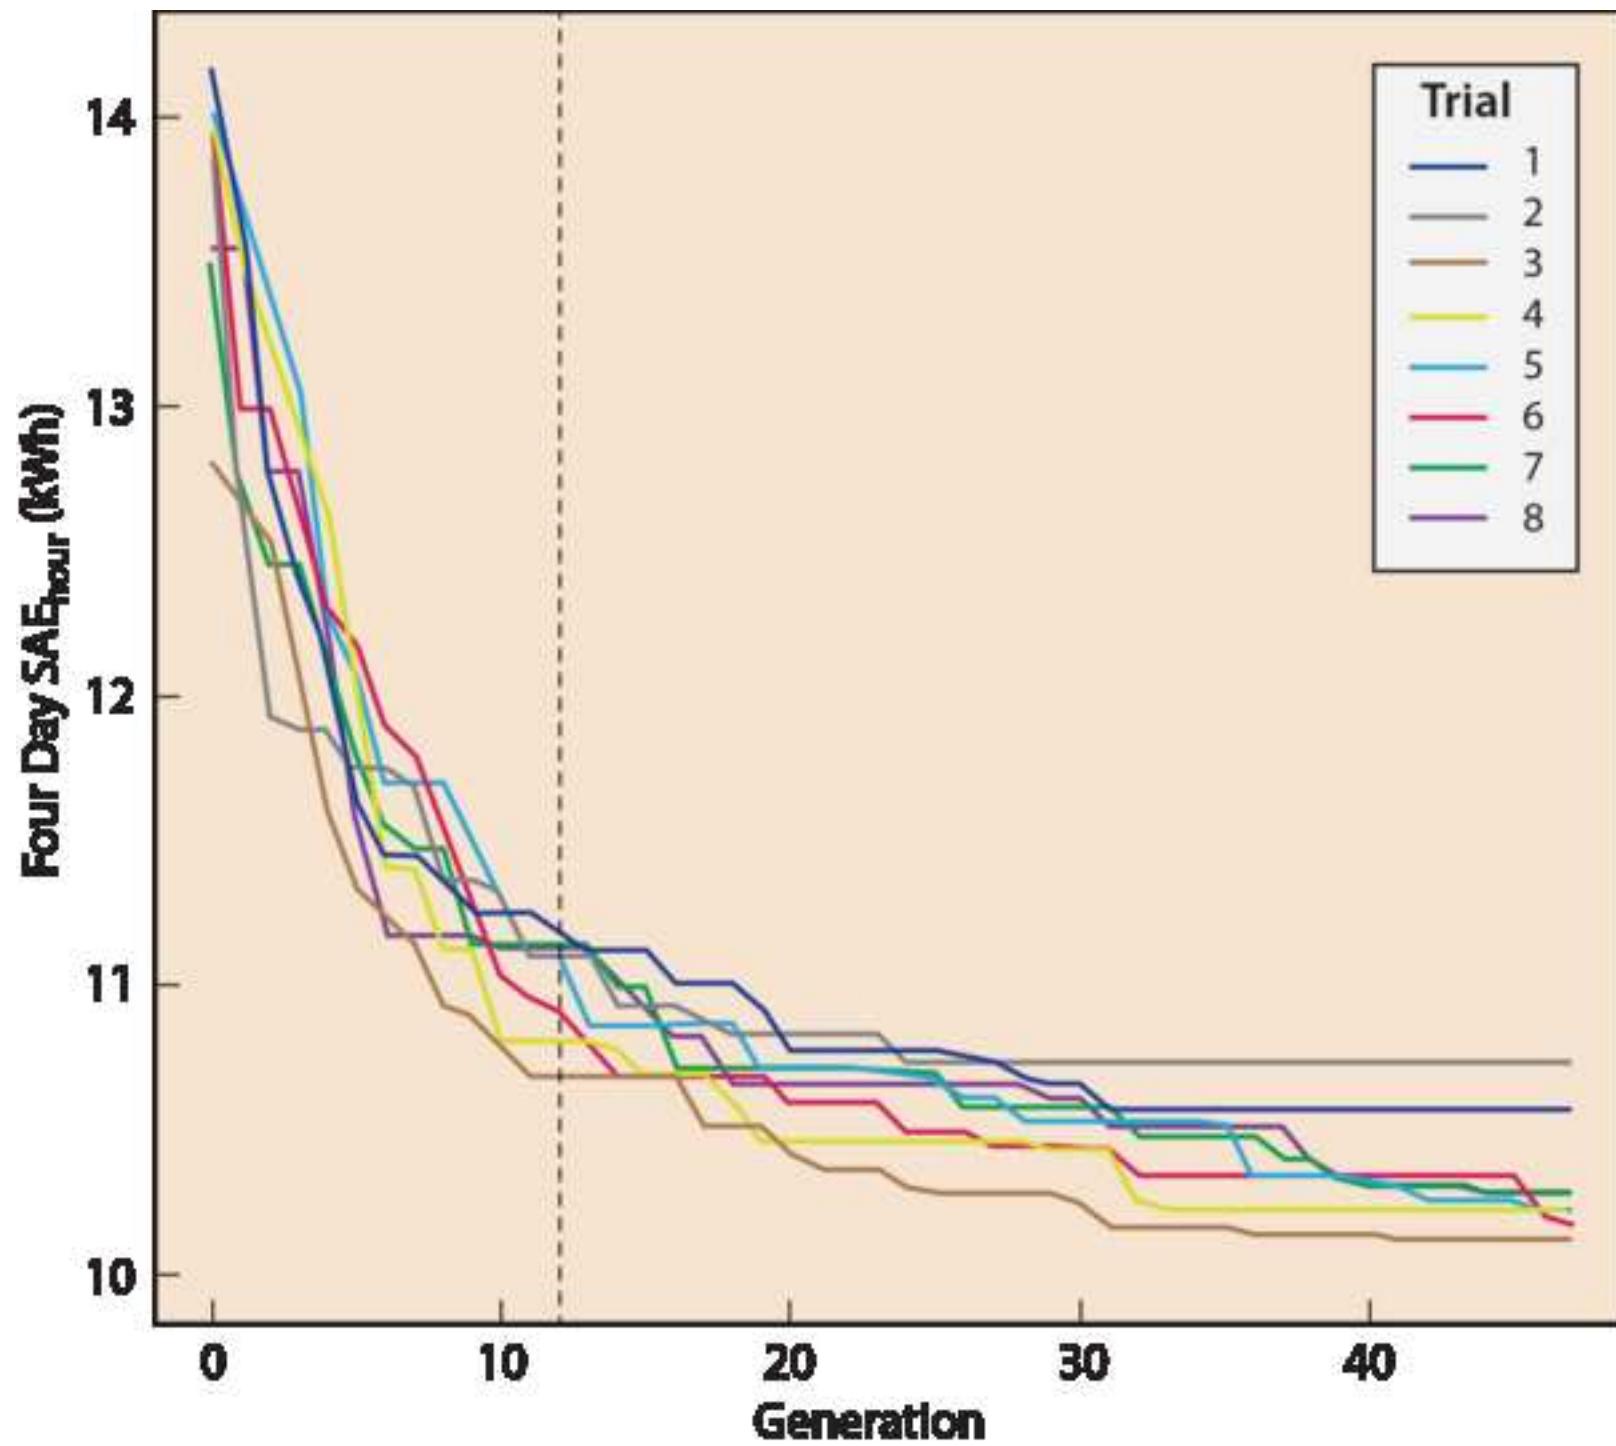




How to Evolve generations of buildings in parallel (\*.psd avail)

[Click here to download high resolution image](#)





## Summary of Experimental Results

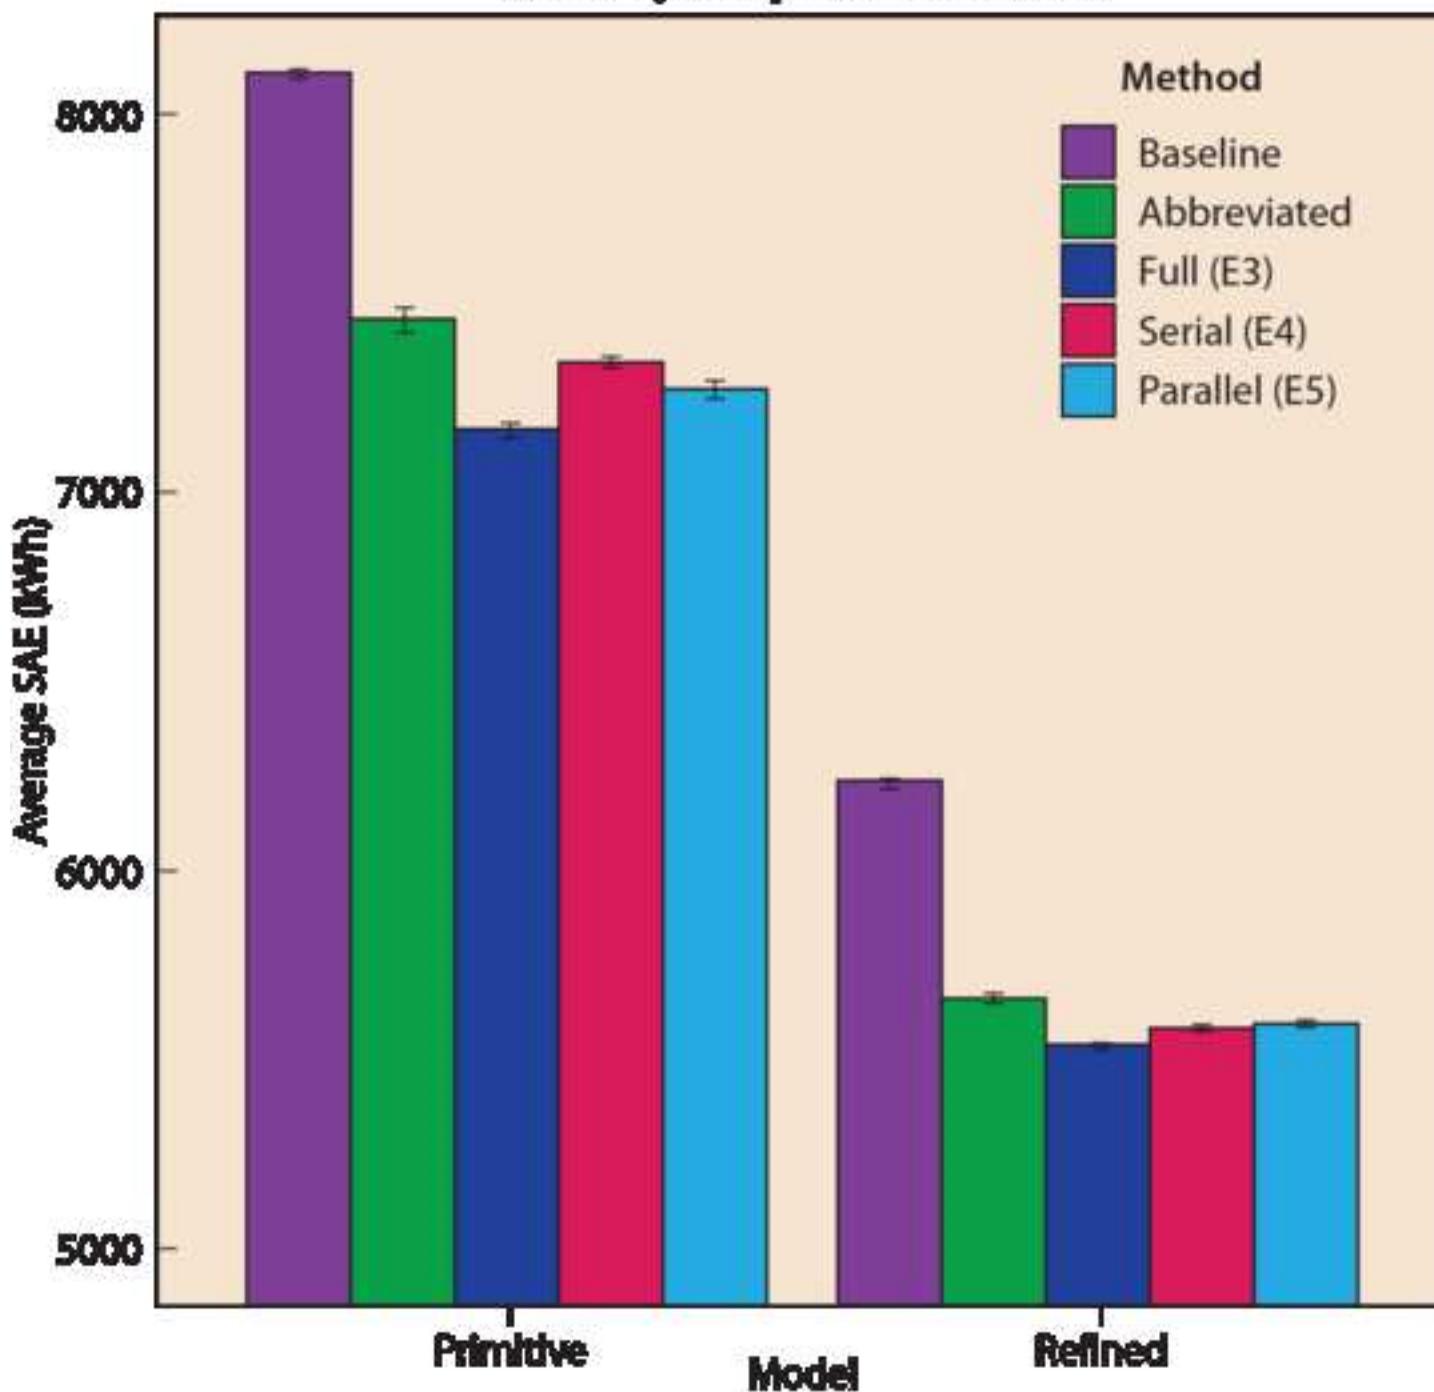


Table 1: The sampling of the SAE to determine correlation was done using an EC with the objective of minimizing the SAE. This produced initial samples with relatively high SAE values and subsequent samples with generally lower SAE as more data were used. The correlation using only the first  $N\%$  of the samples produced shows that even the initial samples are strongly correlated most of the time (with the exception of Trial 3).

Data Used	Trial 1	Trial 2	Trial 3	Trial 4
100%	0.9603	0.9422	0.9015	0.9555
90%	0.8668	0.8438	0.6545	0.8788
80%	0.8683	0.8573	0.6444	0.8922
70%	0.8726	0.8561	0.6231	0.8944
60%	0.8671	0.8489	0.6078	0.8952
50%	0.8703	0.8489	0.5557	0.8971
40%	0.8774	0.8417	0.5663	0.9020
30%	0.8794	0.8208	0.6015	0.8973
20%	0.8841	0.8231	0.5988	0.8860
10%	0.8796	0.8235	0.3985	0.8714

Table 2: In Experiment 2, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the abbreviated schedule. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7539	7578	7683	7638	7495	7467	7503	7587
Refined	Avg	5720	5708	5792	5762	5750	5783	5806	5698
Primitive	Min	7393	7507	7595	7545	7393	7360	7356	7477
Refined	Min	5620	5625	5709	5687	5640	5671	5711	5616

Table 3: In Experiment 3, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the full schedule. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7275	7245	7250	7259	7338	7202	7274	7272
Refined	Avg	5582	5641	5620	5594	5636	5623	5623	5596
Primitive	Min	7175	7138	7150	7122	7253	7114	7162	7179
Refined	Min	5518	5568	5527	5510	5560	5550	5558	5523

Table 4: In Experiment 4, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the abbreviated schedule for the first 75% of the evaluations and the full schedule for the remaining 25%. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7404	7463	7469	7467	7393	7381	7434	7403
Refined	Avg	5626	5670	5688	5614	5695	5667	5642	5646
Primitive	Min	7306	7382	7402	7378	7328	7283	7362	7306
Refined	Min	5564	5586	5608	5545	5596	5618	5566	5563

Table 5: In Experiment 5, an island-model evolutionary computation was used to optimize the primitive and refined models eight independent times (thus 16 different tunings) using a *slow* and a *fast* island. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7340	7330	7435	7470	7391	7349	7390	7356
Refined	Avg	5724	5675	5670	5660	5648	5715	5669	5625
Primitive	Min	7218	7201	7360	7373	7281	7211	7293	7226
Refined	Min	5639	5597	5582	5578	5580	5633	5591	5573

Table 6: The summary comparison of results from Experiments 2–5 show that significant reductions can be gained from tuning with an abbreviated schedule, either alone, serially, or in parallel with full-schedule tuning. Percentages in parentheses are reductions versus the baseline.

Approach	Refined		Primitive	
	SAE (kWh)	RMSE	SAE (kWh)	RMSE
Baseline	6242	1.206	8114	1.625
Abbreviated	5660 (9%)	1.129 (6%)	7453 (8%)	1.514 (7%)
Full	5539 (11%)	1.119 (7%)	7162 (12%)	1.458 (10%)
Serial	5581 (11%)	1.123 (7%)	7343 (9%)	1.497 (8%)
Parallel	5597 (10%)	1.121 (7%)	7270 (10%)	1.482 (9%)

Table 1: The sampling of the SAE to determine correlation was done using an EC with the objective of minimizing the SAE. This produced initial samples with relatively high SAE values and subsequent samples with generally lower SAE as more data were used. The correlation using only the first  $N\%$  of the samples produced shows that even the initial samples are strongly correlated most of the time (with the exception of Trial 3).

Data Used	Trial 1	Trial 2	Trial 3	Trial 4
100%	0.9603	0.9422	0.9015	0.9555
90%	0.8668	0.8438	0.6545	0.8788
80%	0.8683	0.8573	0.6444	0.8922
70%	0.8726	0.8561	0.6231	0.8944
60%	0.8671	0.8489	0.6078	0.8952
50%	0.8703	0.8489	0.5557	0.8971
40%	0.8774	0.8417	0.5663	0.9020
30%	0.8794	0.8208	0.6015	0.8973
20%	0.8841	0.8231	0.5988	0.8860
10%	0.8796	0.8235	0.3985	0.8714

Table 2: In Experiment 2, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the abbreviated schedule. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7539	7578	7683	7638	7495	7467	7503	7587
Refined	Avg	5720	5708	5792	5762	5750	5783	5806	5698
Primitive	Min	7393	7507	7595	7545	7393	7360	7356	7477
Refined	Min	5620	5625	5709	5687	5640	5671	5711	5616

Table 3: In Experiment 3, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the full schedule. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7275	7245	7250	7259	7338	7202	7274	7272
Refined	Avg	5582	5641	5620	5594	5636	5623	5623	5596
Primitive	Min	7175	7138	7150	7122	7253	7114	7162	7179
Refined	Min	5518	5568	5527	5510	5560	5550	5558	5523

Table 4: In Experiment 4, the EC was used to optimize the primitive and refined models eight independent times (thus, 16 different tunings) using the abbreviated schedule for the first 75% of the evaluations and the full schedule for the remaining 25%. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7404	7463	7469	7467	7393	7381	7434	7403
Refined	Avg	5626	5670	5688	5614	5695	5667	5642	5646
Primitive	Min	7306	7382	7402	7378	7328	7283	7362	7306
Refined	Min	5564	5586	5608	5545	5596	5618	5566	5563

Table 5: In Experiment 5, an island-model evolutionary computation was used to optimize the primitive and refined models eight independent times (thus 16 different tunings) using a *slow* and a *fast* island. The final population of 16 individuals for each tuning were used to produce the statistics in this table. The average and minimum  $SAE_{hour}$  (kWh) values for electrical usage from the final populations are shown.

Model	Metric	Trial (kWh)							
		1	2	3	4	5	6	7	8
Primitive	Avg	7340	7330	7435	7470	7391	7349	7390	7356
Refined	Avg	5724	5675	5670	5660	5648	5715	5669	5625
Primitive	Min	7218	7201	7360	7373	7281	7211	7293	7226
Refined	Min	5639	5597	5582	5578	5580	5633	5591	5573

Table 6: The summary comparison of results from Experiments 2–5 show that significant reductions can be gained from tuning with an abbreviated schedule, either alone, serially, or in parallel with full-schedule tuning. Percentages in parentheses are reductions versus the baseline.

Approach	Refined		Primitive	
	SAE (kWh)	RMSE	SAE (kWh)	RMSE
Baseline	6242	1.206	8114	1.625
Abbreviated	5660 (9%)	1.129 (6%)	7453 (8%)	1.514 (7%)
Full	5539 (11%)	1.119 (7%)	7162 (12%)	1.458 (10%)
Serial	5581 (11%)	1.123 (7%)	7343 (9%)	1.497 (8%)
Parallel	5597 (10%)	1.121 (7%)	7270 (10%)	1.482 (9%)

LaTeX Source - 3x images different from attached (similar style)  
[Click here to download LaTeX Source Files: 6\\_Supplementary\\_latex.zip](#)